

Wprowadzenie do sieci neuronowych

Piotr Krzywicki



Przypomnienie

Zadanie klasyfikacji czy regresji polegają na dopasowaniu do otykietowanego zbioru danych X, y funkcji, dla której z dostatecznie dużą dokładnością zachodzi $f_{\theta}(x) \approx y$ dla losowych x, y z tego zbioru.

Przypomnienie

Żeby tego dokonać, uznajemy, że naszą funkcję f_θ daje się sparametryzować wagami θ .

Przypomnienie

Przykładowy schemat wygląda następująco:

0. Do problemu dobierz architekturę modelu, parametryzowaną wagami θ , funkcję błędu $\mathcal{L}(x, y, \theta)$, której niska oczekiwana wartość dla danych wag gwarantuje dobrą jakość modelu
1. Niech X to dane wejściowe, y to etykiety
2. Zainicjalizuj parametry modelu θ
3. Powtarzaj dopóki $\mathcal{L}(\theta)$ nie będzie wystarczająco niska
 - Minimalizuj $\mathcal{L}(\theta)$ zmieniając wagi θ
 - zazwyczaj powyższa minimalizacja oparta jest na gradiencie
 - Innymi słowy zazwyczaj $\theta = \theta - \alpha * \nabla_{\theta} \mathcal{L}(\theta)$

Przypomnienie

Algorytm regresji liniowej zakłada:

$$f_{\theta}(x) = \theta^T x$$
$$\mathcal{L}(x, y, \theta) = (f_{\theta}(x) - y)^2$$

Algorytm regresji logistycznej zakłada:

$$f_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$
$$\mathcal{L}(x, y, \theta) = H(f_{\theta}(x), y)$$

Algorytm regresji logistycznej z wieloma klasami (softmax regression/multinomial regression) zakłada:

$$f_{\theta}(x) = \text{softmax}(\theta^T x)$$

gdzie

$$\text{softmax}(x_1, \dots, x_n) = \left(\frac{\exp(x_i)}{\exp(x_1) + \dots + \exp(x_n)} \right)_i$$
$$\mathcal{L}(x, y, \theta) = H(f_{\theta}(x), y)$$

DATA

Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+


-

0 HIDDEN LAYERS


OUTPUT

Test loss 0.000

Training loss 0.000



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0




Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?




+ - 0 HIDDEN LAYERS


OUTPUT

Test loss 0.540

Training loss 0.497



Colors shows data, neuron and weight values.



☐ Show test data ☒ Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0




Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+


-

0 HIDDEN LAYERS


OUTPUT

Test loss 0.512

Training loss 0.501



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES


Which properties do you want to feed in?

- X_1
- X_2
- X_1^2
- X_2^2
- $X_1 X_2$
- $\sin(X_1)$
- $\sin(X_2)$

+ - 0 HIDDEN LAYERS

OUTPUT

Test loss 0.486
Training loss 0.461




Colors shows data, neuron and weight values.

☐ Show test data ☒ Discretize output

Przypomnienie – feature engineering

Zwizualizujemy dopasowanie naszego modelu:


```
[ ] xs = np.linspace(-3, 3, 1000)
    plt.scatter(X[:, 0], y)
    plt.scatter(xs,
        """
        TODO:
        your code goes here
        """
    )
    plt.show()
```




Jak rozwiązać te problemy?

- Feature engineering (inżynieria cech)
- Silniejszy, nieliniowy model

Model Regresji Liniowej, ale inaczej




Dlaczego by nie dołożyć dodatkowych węzłów?




Nazewnictwo


'input' layer



Nazewnictwo



Nazewnictwo



Problem? To nie jest silniejsza klasa modeli

$$y = W_2 W_1 x = W' x, \text{ gdzie } W = W_2 W_1$$

Złamanie liniowości

$$y = f(W_2 f(W_1 x)), \text{ gdzie } f \text{ nieliniowa}$$

Sprawdzone wybory funkcji aktywacji f

Rectified
Linear Unit
(ReLU)



$$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$




$$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$$

Logistic (a.k.a
Soft step)



$$f(x) = \frac{1}{1 + e^{-x}}$$

ArcTan



$$f(x) = \tan^{-1}(x)$$

DATA

Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES


Which properties do you want to feed in?



+ - 1 HIDDEN LAYER

+ -


3 neurons




OUTPUT

Test loss 0.007

Training loss 0.004



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

- X_1
- X_2
- X_1^2
- X_2^2
- $X_1 X_2$
- $\sin(X_1)$
- $\sin(X_2)$

+ - 1 HIDDEN LAYER


+ -

2 neurons

This is the output from one *neuron*. Hover to see it larger.

OUTPUT

Test loss 0.243
Training loss 0.246



Colors shows data, neuron and weight values.

☐ Show test data ☒ Discretize output

DATA

Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+ - 1 HIDDEN LAYER

+ -


2 neurons

This is the output from one **neuron**.
Hover to see it larger.

OUTPUT

Test loss 0.181

Training loss 0.159




Colors shows data, neuron and weight values.




☐ Show test data

☒ Discretize output


3 hidden neurons




6 hidden neurons




20 hidden neurons




Przypomnienie – problem przeuczenia (overfitting)




$\lambda = 0.001$

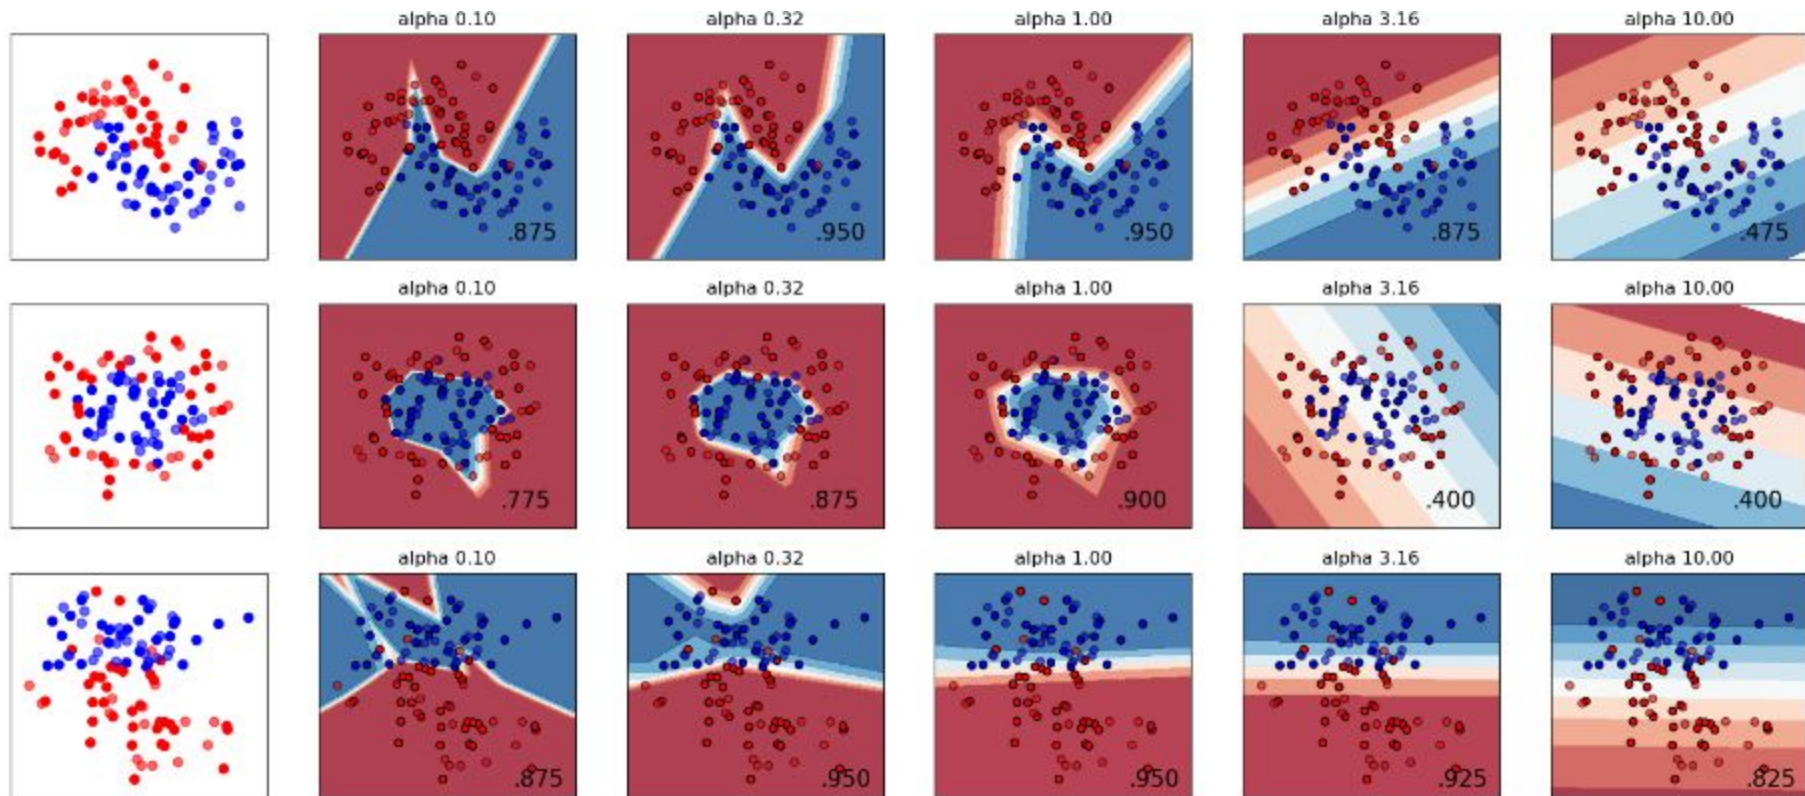


$\lambda = 0.01$

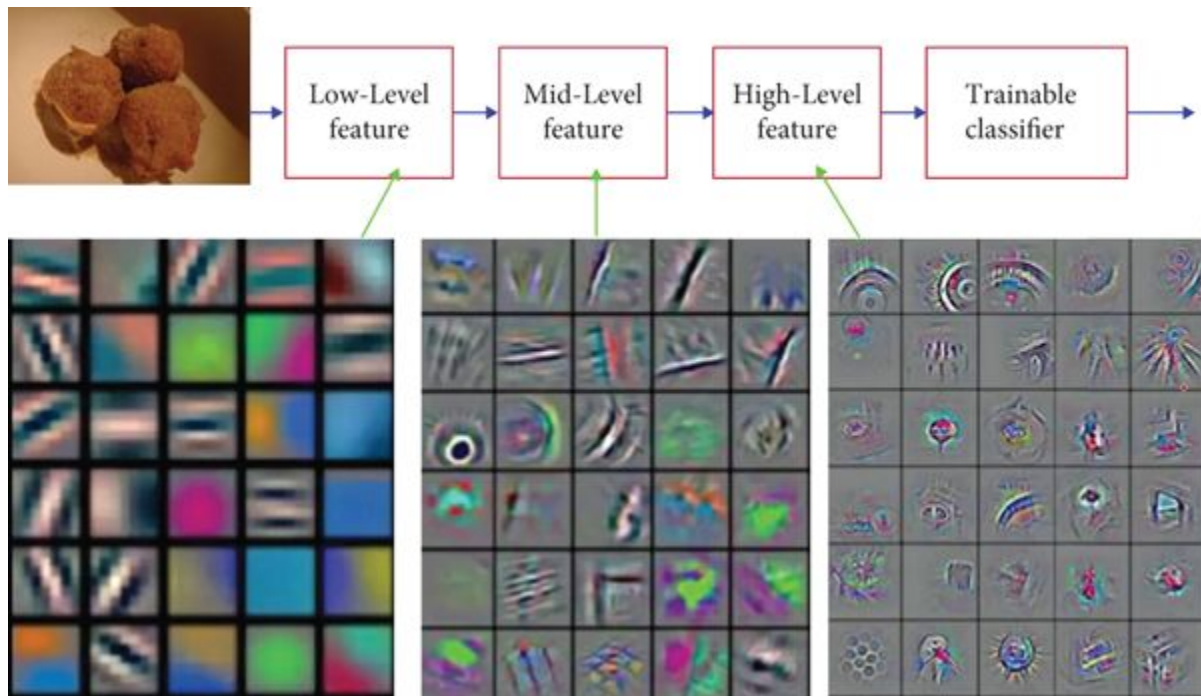


$\lambda = 0.1$





Hierarchiczność/kompozycyjność reprezentacji



Zadanie labowe 0



- Przykłady na slajdach zostały stworzone z użyciem <https://playground.tensorflow.org/>
- Proszę poeksperymentować ze:
 - Zbiorami danych
 - Architekturą sieci neuronowej
 - Współczynnikiem uczenia (learning rate)
 - Funkcją aktywacji


Zadanie labowe 1

- Będziemy rozwiązywać zadanie klasyfikacji z użyciem sieci neuronowych
- Będziemy pracować na zbiorze danych CIFAR10
- Porównamy jakość modelu sieci neuronowej z modelem liniowym

Zadanie labowe 1

Notebook settings

Hardware accelerator
GPU  

GPU class
Standard 

Want access to premium GPUs? [Purchase additional compute units](#)

☐ Omit code cell output when saving this notebook

Cancel [Save](#)

Zadanie labowe 1

```
▶ !nvidia-smi

Thu Apr 13 07:00:44 2023

+-----+
| NVIDIA-SMI 525.85.12      Driver Version: 525.85.12      CUDA Version: 12.0      |
+-----+-----+-----+-----+-----+-----+
| GPU  Name          Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp   Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
|====+=====+====+=====+=====+=====+
|  0 Tesla T4             Off      | 00000000:00:04.0 Off |             0        |
| N/A   71C    P8       11W /  70W |  0MiB / 15360MiB |      0%      Default |
+-----+-----+-----+-----+-----+-----+

+-----+
| Processes: |
| GPU   GI    CI          PID    Type    Process name                        GPU Memory |
|      ID    ID                                   |             Usage |
+-----+-----+-----+-----+-----+-----+
| No running processes found |
+-----+
```

Zadanie labowe 1

Krok 1: Import podstawowy bibliotek

```
[2] import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
```


Zadanie labowe 1

▼ Krok 2: Wczytanie zbioru danych CIFAR10

✓
9s



```
# Define the transforms to be applied to the CIFAR-10 dataset
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
```

Zadanie labowe 1

```
# Load the CIFAR-10 training and test datasets
trainset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transform
)
```

Zadanie labowe 1

```
trainloader = torch.utils.data.DataLoader(  
    trainset,  
    batch_size=4,  
    shuffle=True,  
    num_workers=2  
)
```

Zadanie labowe 1

```
testset = torchvision.datasets.CIFAR10(  
    root='./data',  
    train=False,  
    download=True,  
    transform=transform  
)  
  
testloader = torch.utils.data.DataLoader(  
    testset,  
    batch_size=4,  
    shuffle=False,  
    num_workers=2  
)
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz
100%|██████████| 170498071/170498071 [00:03<00:00, 43232169.68it/s]
Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

Zadanie labowe 1

Krok 3: Wizualizacja zbioru danych CIFAR10

```
import matplotlib.pyplot as plt
import numpy as np

# Define the classes for the CIFAR-10 dataset
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')


# Get some random training images
dataiter = iter(trainloader)
images, labels = dataiter.__next__()

# Show the images
fig, axes = plt.subplots(1, len(images), figsize=(10,2))
for i, image in enumerate(images):
    # Unnormalize the image
    image = image / 2 + 0.5
    np_image = image.numpy()

    # Transpose the channels to display the image
    transposed = np.transpose(np_image, (1, 2, 0))

    # Show the image
    axes[i].imshow(transposed)
    axes[i].set_title(classes[labels[i]])

# Display the images
plt.show()
```



Zadanie labowe 1

Krok 4. Definicja modelu sieci neuronowej oraz modelu liniowego

```
class MLP(nn.Module):  
    def __init__(self):  
        super(MLP, self).__init__()  
        self.fc1 = nn.Linear(32*32*3, 512)  
        self.fc2 = nn.Linear(512, 256)  
        self.fc3 = nn.Linear(256, 10)  
  
    def forward(self, x):  
        """  
        TODO: your code goes here.  
        Hints:  
        * flatten the image before passing it through layers  
        * use self.fc1, self.fc2, self.fc3  
        * apply nonlinearity nn.functional.relu after every layer  
        """  
        return x
```

Zadanie labowe 1

```
class LinearModel(nn.Module):
    def __init__(self):
        super(LinearModel, self).__init__()
        self.linear = nn.Linear([32*32*3, 10])

    def forward(self, x):
        """
        TODO: your code goes here.
        Hints:
        * flatten the image before passing it through layers
        * use self.linear
        """
        return x
```

Zadanie labowe 1

Krok 5. Implementacja pętli uczącej

```
[ ] def train(net, criterion, optimizer, trainloader, device):  
    running_loss = 0.0  
    for i, data in enumerate(trainloader, 0):  
        inputs, labels = data  
        """  
        TODO: move the inputs and labels to device  
        """  
  
        optimizer.zero_grad()  
  
        """  
        TODO:  
        1. pass the inputs through the model to obtain the outputs  
        2. calculate the criterion loss, based on the outputs and labels  
        3. calculate the gradients using loss object  
        4. perform the optimization step using optimizer  
        """  
  
        running_loss += loss.item()  
  
    return running_loss / len(trainloader)
```


Zadanie labowe 1

Krok 6. Wyuczenie modeli

```
▶ # Define the device to be used
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

# Define the MLP model and optimizer
mlp_net = MLP().to(device)
mlp_criterion = nn.CrossEntropyLoss()
mlp_optimizer = optim.SGD(mlp_net.parameters(), lr=0.001, momentum=0.9)
"""
TODO:
    experiment with the learning rate
"""
```

Zadanie labowe 1

```
# Define the linear model and optimizer
linear_net = LinearModel().to(device)
linear_criterion = nn.CrossEntropyLoss()
linear_optimizer = optim.SGD(linear_net.parameters(), lr=0.001, momentum=0.9)
"""
    TODO:
        experiment with the learning rate
"""
```

Zadanie labowe 1

```
# Define the number of epochs
epochs = 10
"""
    TODO:
        experiment with length of the training, does adding more epochs helps?
"""

# Train the MLP model
for epoch in range(epochs):
    train_mlp_loss = train(mlp_net, mlp_criterion, mlp_optimizer, trainloader, device)
    print('Epoch %d MLP loss: %.3f' % (epoch + 1, train_mlp_loss))

# Train the linear model
for epoch in range(epochs):
    train_linear_loss = train(linear_net, linear_criterion, linear_optimizer, trainloader, device)
    print('Epoch %d Linear loss: %.3f' % (epoch + 1, train_linear_loss))
```

Zadanie labowe 1

```
Epoch 1 MLP loss: 1.649
Epoch 2 MLP loss: 1.419
Epoch 3 MLP loss: 1.304
Epoch 4 MLP loss: 1.215
Epoch 5 MLP loss: 1.136
Epoch 6 MLP loss: 1.066
Epoch 7 MLP loss: 0.993
Epoch 8 MLP loss: 0.930
Epoch 9 MLP loss: 0.870
Epoch 10 MLP loss: 0.810
Epoch 1 Linear loss: 2.153
Epoch 2 Linear loss: 2.106
Epoch 3 Linear loss: 2.080
Epoch 4 Linear loss: 2.067
Epoch 5 Linear loss: 2.057
Epoch 6 Linear loss: 2.057
Epoch 7 Linear loss: 2.046
Epoch 8 Linear loss: 2.046
Epoch 9 Linear loss: 2.039
Epoch 10 Linear loss: 2.030
```

Zadanie labowe 1

Krok 7. Ewaluacja modeli na zbiorze testowym

```
▶ # Evaluate the models on the test set
mlp_correct = 0
linear_correct = 0
total = 0
with torch.no_grad():
    for data in testloader:
        images, labels = data
        images, labels = images.to(device), labels.to(device)
        mlp_outputs = mlp_net(images)
        _, mlp_predicted = torch.max(mlp_outputs.data, 1)
        linear_outputs = linear_net(images)
        _, linear_predicted = torch.max(linear_outputs.data, 1)
        total += labels.size(0)
        mlp_correct += (mlp_predicted == labels).sum().item()
        linear_correct += (linear_predicted == labels).sum().item()

mlp_accuracy = "" TODO: your code goes here ""
linear_accuracy = "" TODO: your code goes here ""

print('MLP Accuracy on the test images: %d %%' % mlp_accuracy)
print('Linear Accuracy on the test images: %d %%' % linear_accuracy)
```

```
➡ MLP Accuracy on the test images: 53 %
   Linear Accuracy on the test images: 33 %
```

Zadanie labowe 1

Krok 8. Wizualizacja predykcji obu modeli

```
# Get a batch of test images
testloader = torch.utils.data.DataLoader(
    testset,
    batch_size=32,
    shuffle=False,
    num_workers=2
)

dataiter = iter(testloader)
images, labels = dataiter.__next__()

# Move the images to the device
images = images.to(device)

# Get the MLP predictions and output probabilities
mlp_outputs = mlp_net(images)
mlp_predicted = torch.max(mlp_outputs, 1)[1]
mlp_probs = torch.nn.functional.softmax(mlp_outputs, dim=1)

# Get the linear predictions and output probabilities
linear_outputs = linear_net(images.view(images.size(0), -1))
linear_predicted = torch.max(linear_outputs, 1)[1]
linear_probs = torch.nn.functional.softmax(linear_outputs, dim=1)
```

Zadanie labowe 1

```
# Define the labels for the classes
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck')

# Define a function to show an image along with its predicted class and probabilities
def imshow_probs(img, title, probs):
    img = img / 2 + 0.5      # unnormalize the image
    npimg = img.numpy()
    fig, (ax1, ax2) = plt.subplots(ncols=2, figsize=(8,3))
    ax1.imshow(np.transpose(npimg, (1, 2, 0)))
    ax1.set_title(title)
    ax1.axis('off')
    ax2.bar(classes, probs)
    ax2.set_xticklabels(classes, rotation=45)
    ax2.set_ylim([0, 1])
    ax2.set_ylabel('Probability')
    ax2.set_title('Class Probabilities')
    fig.tight_layout()

# Show the images with their predicted classes and probabilities for the MLP model
for i in range(10):
    print("MLP:")
    mlp_probs_i = mlp_probs[i].cpu().detach().numpy()
    imshow_probs(images[i].cpu(), classes[mlp_predicted[i]], mlp_probs_i)
    plt.show()
    print("Linear:")
    linear_probs_i = linear_probs[i].cpu().detach().numpy()
    imshow_probs(images[i].cpu(), classes[linear_predicted[i]], linear_probs_i)
    plt.show()
```

Zadanie labowe 1


MLP:

```
<ipython-input-11-f1c0206d2b07>:37: UserWarning: FixedFormatter should only be used together with FixedLocator  
ax2.set_xticklabels(classes, rotation=45)
```

cat



Class Probabilities




Linear:

cat




Class Probabilities



Zadanie labowe 1


MLP:

ship



Linear:


ship



Zadanie labowe 1


MLP:

truck



Linear:


ship



Zadanie labowe 1


MLP:

horse



Linear:


plane



Zadanie labowe 1


MLP:

deer



Linear:


deer



Zadanie labowe 1


MLP:

frog



Linear:


cat



Zadanie labowe 1


MLP:

car



Linear:


cat



Zadanie labowe 1


MLP:

frog



Linear:


bird



Zadanie labowe 1


MLP:

dog



Linear:


dog



Zadanie labowe 1

MLP:

car



Linear:

car

